

FINAL PROJECT SUBMISSION

Student name: AUGUSTINE MAGANI

Student pace: self paced / part time / full time:FULL TIME

Instructor name: FAITH ROTICH

FILM PRODUCTION ANALYSIS

INTRO

Data-driven decision-making has become crucial to ensuring successful investments, targeted marketing, and sustained growth in the film production industry.

This analysis focuses on evaluating and comparing the performance of films listed on several key platforms i.e; Box Office Mojo, The Numbers, and The Movie Database (TMDb).

BUSINESS UNDERSTANDING

A production company must understand what factors contribute to a movie's success commercially before investing in new films. This entails revenue and also how movies perform across audience platforms.

Business Questions

Which movie language performs best among different platforms

What trends exist across genres, release periods, and production budgets in relation to success?

What is the ideal runtime minutes for a movie?

Can early popularity forecast long-term success?

Objectives

Identifying the top ranking movie genres over time

Examining trends in average gross revenue of films over time.

Determining the release months with the highest audience engagement

Examining movie production budget over the years

Identifying ROI trends over the years

Ranking the top 10 highest-grossing film studios

Determining the most popular movie languages

DATA UNDERSTANDING

For this analysis, I used 3 datasets.

Box Office Mojo

Providing revenue statistics for movies published in the Website.

The Movie Database(TMDb)

Also provides useful elements for movies published in the Movie database

The Numbers

Provides revenue data from The numbers dataset

DATA PREPARATION

In [6]: *#Importing all the necessary libraries*

```
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
from numbers import Number
import sqlite3
from scipy import stats
%matplotlib inline
import os
import zipfile
```

Loading the Box Office Mojo dataset

In [7]: `bom_df = pd.read_csv("bom.movie_gross.csv.gz")`

Reading the first 5 rows

In [8]: `bom_df.head()`

Out[8]:

| | title | studio | domestic_gross | foreign_gross | year |
|---|---|--------|----------------|---------------|------|
| 0 | Toy Story 3 | BV | 415000000.0 | 652000000 | 2010 |
| 1 | Alice in Wonderland (2010) | BV | 334200000.0 | 691300000 | 2010 |
| 2 | Harry Potter and the Deathly Hallows Part 1 | WB | 296000000.0 | 664300000 | 2010 |
| 3 | Inception | WB | 292600000.0 | 535700000 | 2010 |
| 4 | Shrek Forever After | P/DW | 238700000.0 | 513900000 | 2010 |

Viewing the last rows of the df

In [9]: `bom_df.tail()`

Out[9]:

| | title | studio | domestic_gross | foreign_gross | year |
|------|-----------------------------|------------|----------------|---------------|------|
| 3382 | The Quake | Magn. | 6200.0 | NaN | 2018 |
| 3383 | Edward II (2018 re-release) | FM | 4800.0 | NaN | 2018 |
| 3384 | El Pacto | Sony | 2500.0 | NaN | 2018 |
| 3385 | The Swan | Synergetic | 2400.0 | NaN | 2018 |
| 3386 | An Actor Prepares | Grav. | 1700.0 | NaN | 2018 |

Identifying the number of rows and columns

In [34]: `bom_df.shape`

Out[34]: (3387, 5)

The box office dataframe has 3,387 rows and 5 columns.

In [33]: `bom_df.columns`

Out[33]: Index(['title', 'studio', 'domestic_gross', 'foreign_gross', 'year'], dtype='object')

title Contains the title of the movies

studio column has abbreviated names of the studios producing the respective movie

domestic_gross column shows the amount of revenue generated by the movie through sales, locally

foreign_gross shows the amount of revenue generated by the movie internationally

year shows the year when the movie was premiered

Showing a summary of the df

In [32]: bom_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   title                 3387 non-null   object
1   studio                3382 non-null   object
2   domestic_gross        3359 non-null   float64
3   foreign_gross         2037 non-null   object
4   year                  3387 non-null   int64
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB
```

In [31]: *#range of the rows of the df*
bom_df.index

Out[31]: RangeIndex(start=0, stop=3387, step=1)

In [37]: bom_df.dtypes

```
Out[37]: title                object
studio                  object
domestic_gross         float64
foreign_gross           object
year                   int64
dtype: object
```

Detecting NaN Values

In [39]: *#All cells containing NaN are converted to True, and all cells containing valid data are convert*
bom_df.isna()

```
Out[39]:
```

| | title | studio | domestic_gross | foreign_gross | year |
|------|-------|--------|----------------|---------------|-------|
| 0 | False | False | False | False | False |
| 1 | False | False | False | False | False |
| 2 | False | False | False | False | False |
| 3 | False | False | False | False | False |
| 4 | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... |
| 3382 | False | False | False | True | False |
| 3383 | False | False | False | True | False |
| 3384 | False | False | False | True | False |
| 3385 | False | False | False | True | False |
| 3386 | False | False | False | True | False |

3387 rows × 5 columns

Number of NaN values in the dataset by columns

In [40]: bom_df.isna().sum()

```
Out[40]: title                0
studio                  5
domestic_gross         28
foreign_gross        1350
year                   0
dtype: int64
```

Filling NaN with no values

In [41]: `bom_df.fillna('')`

Out[41]:

| | title | studio | domestic_gross | foreign_gross | year |
|------|---|------------|----------------|---------------|------|
| 0 | Toy Story 3 | BV | 415000000.0 | 652000000 | 2010 |
| 1 | Alice in Wonderland (2010) | BV | 334200000.0 | 691300000 | 2010 |
| 2 | Harry Potter and the Deathly Hallows Part 1 | WB | 296000000.0 | 664300000 | 2010 |
| 3 | Inception | WB | 292600000.0 | 535700000 | 2010 |
| 4 | Shrek Forever After | P/DW | 238700000.0 | 513900000 | 2010 |
| ... | ... | ... | ... | ... | ... |
| 3382 | The Quake | Magn. | 6200.0 | | 2018 |
| 3383 | Edward II (2018 re-release) | FM | 4800.0 | | 2018 |
| 3384 | El Pacto | Sony | 2500.0 | | 2018 |
| 3385 | The Swan | Synergetic | 2400.0 | | 2018 |
| 3386 | An Actor Prepares | Grav. | 1700.0 | | 2018 |

3387 rows × 5 columns

THE MOVIE DATABASE

Loading the tm dataset

In [34]: `tmdb_df = pd.read_csv("tmdb.movies.csv.gz")`

In [35]: `#Viewing the first five rows`
`tmdb_df.head()`

Out[35]:

| | Unnamed: 0 | genre_ids | id | original_language | original_title | popularity | release_date | title | vote_average | vote_ |
|---|------------|---------------------|-------|-------------------|--|------------|--------------|--|--------------|-------|
| 0 | 0 | [12, 14, 10751] | 12444 | en | Harry Potter and the Deathly Hallows: Part 1 | 33.533 | 2010-11-19 | Harry Potter and the Deathly Hallows: Part 1 | 7.7 | |
| 1 | 1 | [14, 12, 16, 10751] | 10191 | en | How to Train Your Dragon | 28.734 | 2010-03-26 | How to Train Your Dragon | 7.7 | |
| 2 | 2 | [12, 28, 878] | 10138 | en | Iron Man 2 | 28.515 | 2010-05-07 | Iron Man 2 | 6.8 | |
| 3 | 3 | [16, 35, 10751] | 862 | en | Toy Story | 28.005 | 1995-11-22 | Toy Story | 7.9 | |
| 4 | 4 | [28, 878, 12] | 27205 | en | Inception | 27.920 | 2010-07-16 | Inception | 8.3 | |

```
In [36]: #Checking the last five rows
tmdb_df.tail()
```

```
Out[36]:
```

| | Unnamed: 0 | genre_ids | id | original_language | original_title | popularity | release_date | title | v |
|-------|------------|-----------------|--------|-------------------|-----------------------|------------|--------------|-----------------------|---|
| 26512 | 26512 | [27, 18] | 488143 | en | Laboratory Conditions | 0.6 | 2018-10-13 | Laboratory Conditions | |
| 26513 | 26513 | [18, 53] | 485975 | en | _EXHIBIT_84xxx_ | 0.6 | 2018-05-01 | _EXHIBIT_84xxx_ | |
| 26514 | 26514 | [14, 28, 12] | 381231 | en | The Last One | 0.6 | 2018-10-01 | The Last One | |
| 26515 | 26515 | [10751, 12, 28] | 366854 | en | Trailer Made | 0.6 | 2018-06-22 | Trailer Made | |
| 26516 | 26516 | [53, 27] | 309885 | en | The Church | 0.6 | 2018-10-05 | The Church | |

```
In [37]: #Identifying the number of rows ad columns
tmdb_df.shape
```

```
Out[37]: (26517, 10)
```

The movie database has 26,517 rows and 10 columns

```
In [38]: #Showing the columns
tmdb_df.columns
```

```
Out[38]: Index(['Unnamed: 0', 'genre_ids', 'id', 'original_language', 'original_title',
               'popularity', 'release_date', 'title', 'vote_average', 'vote_count'],
              dtype='object')
```

```
In [39]: #Showing a summary of the tmdb df
tmdb_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26517 entries, 0 to 26516
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            26517 non-null  int64
1   genre_ids             26517 non-null  object
2   id                    26517 non-null  int64
3   original_language     26517 non-null  object
4   original_title        26517 non-null  object
5   popularity            26517 non-null  float64
6   release_date          26517 non-null  object
7   title                 26517 non-null  object
8   vote_average          26517 non-null  float64
9   vote_count            26517 non-null  int64
dtypes: float64(2), int64(3), object(5)
memory usage: 2.0+ MB
```

```
In [ ]: Detecting the NaN values
```

Number of NaN values in the dataset by columns

```
In [40]: tmdb_df.isna().sum()
```

```
Out[40]: Unnamed: 0      0
genre_ids      0
id             0
original_language  0
original_title  0
popularity     0
release_date   0
title          0
vote_average   0
vote_count     0
dtype: int64
```

Now it is evident there are no NaN values in the movies database thus the data is clean.

Checking for missing/null values in the dataset

```
In [41]: tmdb_df.isnull()
```

```
Out[41]:
```

| | Unnamed: 0 | genre_ids | id | original_language | original_title | popularity | release_date | title | vote_average | vote |
|-------|------------|-----------|-------|-------------------|----------------|------------|--------------|-------|--------------|-------|
| 0 | False | False | False | False | False | False | False | False | False | False |
| 1 | False | False | False | False | False | False | False | False | False | False |
| 2 | False | False | False | False | False | False | False | False | False | False |
| 3 | False | False | False | False | False | False | False | False | False | False |
| 4 | False | False | False | False | False | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 26512 | False | False | False | False | False | False | False | False | False | False |
| 26513 | False | False | False | False | False | False | False | False | False | False |
| 26514 | False | False | False | False | False | False | False | False | False | False |
| 26515 | False | False | False | False | False | False | False | False | False | False |
| 26516 | False | False | False | False | False | False | False | False | False | False |

26517 rows × 10 columns

Counting missing values per column

```
In [42]: tmdb_df.isnull().sum()
```

```
Out[42]: Unnamed: 0      0
genre_ids      0
id             0
original_language  0
original_title  0
popularity     0
release_date   0
title          0
vote_average   0
vote_count     0
dtype: int64
```

As you can see, this dataset has not a single missing value and is intact

THE NUMBERS DATASET

```
In [17]: #Loading the dataset
tn_budget = pd.read_csv("tn.movie_budgets.csv.gz")
```

```
In [18]: #Viewing the first five rows
tn_budget.head()
```

```
Out[18]:
```

| | id | release_date | movie | production_budget | domestic_gross | worldwide_gross |
|---|----|--------------|---|-------------------|----------------|-----------------|
| 0 | 1 | Dec 18, 2009 | Avatar | \$425,000,000 | \$760,507,625 | \$2,776,345,279 |
| 1 | 2 | May 20, 2011 | Pirates of the Caribbean: On Stranger Tides | \$410,600,000 | \$241,063,875 | \$1,045,663,875 |
| 2 | 3 | Jun 7, 2019 | Dark Phoenix | \$350,000,000 | \$42,762,350 | \$149,762,350 |
| 3 | 4 | May 1, 2015 | Avengers: Age of Ultron | \$330,600,000 | \$459,005,868 | \$1,403,013,963 |
| 4 | 5 | Dec 15, 2017 | Star Wars Ep. VIII: The Last Jedi | \$317,000,000 | \$620,181,382 | \$1,316,721,747 |

```
In [19]: #Viewing the last five rows
tn_budget.tail()
```

```
Out[19]:
```

| | id | release_date | movie | production_budget | domestic_gross | worldwide_gross |
|------|----|--------------|-------------------------------|-------------------|----------------|-----------------|
| 5777 | 78 | Dec 31, 2018 | Red 11 | \$7,000 | \$0 | \$0 |
| 5778 | 79 | Apr 2, 1999 | Following | \$6,000 | \$48,482 | \$240,495 |
| 5779 | 80 | Jul 13, 2005 | Return to the Land of Wonders | \$5,000 | \$1,338 | \$1,338 |
| 5780 | 81 | Sep 29, 2015 | A Plague So Pleasant | \$1,400 | \$0 | \$0 |
| 5781 | 82 | Aug 5, 2005 | My Date With Drew | \$1,100 | \$181,041 | \$181,041 |

DATA CLEANING

Box Office Mojo

Detecting NaN values

```
In [93]: bom_df.isna()
```

```
Out[93]:
```

| | title | studio | domestic_gross | foreign_gross | year |
|------|-------|--------|----------------|---------------|-------|
| 0 | False | False | False | False | False |
| 1 | False | False | False | False | False |
| 2 | False | False | False | False | False |
| 3 | False | False | False | False | False |
| 4 | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... |
| 3382 | False | False | False | True | False |
| 3383 | False | False | False | True | False |
| 3384 | False | False | False | True | False |
| 3385 | False | False | False | True | False |
| 3386 | False | False | False | True | False |

3387 rows × 5 columns

All cells containing NaN are converted to True, and all cells containing valid data are converted to False

```
In [92]: #Counting missing values per column

round((bom_df.isnull().sum()/ len(bom_df))*100,2)
```

```
Out[92]: title          0.00
studio          0.15
domestic_gross  0.83
foreign_gross   39.86
year            0.00
dtype: float64
```

foreign_gross holds the highest number of missing values, with 39.86% of its data missing. This can be calculated using other columns like domestic_gross and creating another column called total_gross then computing using the mean of the total gross column.

```
In [94]: # Convert the two numeric number columns from strings to float datatypes

bom_df['foreign_gross'] = pd.to_numeric(bom_df['foreign_gross'], errors='coerce')

# Calculate another column Total_gross for total gross generated both locally and internationally

bom_df['Total_gross'] = bom_df['domestic_gross'] + bom_df['foreign_gross']
bom_df.head()
```

```
Out[94]:
```

| | title | studio | domestic_gross | foreign_gross | year | Total_gross |
|---|---|--------|----------------|---------------|------|--------------|
| 0 | Toy Story 3 | BV | 415000000.0 | 652000000.0 | 2010 | 1.067000e+09 |
| 1 | Alice in Wonderland (2010) | BV | 334200000.0 | 691300000.0 | 2010 | 1.025500e+09 |
| 2 | Harry Potter and the Deathly Hallows Part 1 | WB | 296000000.0 | 664300000.0 | 2010 | 9.603000e+08 |
| 3 | Inception | WB | 292600000.0 | 535700000.0 | 2010 | 8.283000e+08 |
| 4 | Shrek Forever After | P/DW | 238700000.0 | 513900000.0 | 2010 | 7.526000e+08 |

```
In [96]: # Use the mean total gross to fill missing values for foreign gross

mean_total = bom_df['Total_gross'].mean()
bom_df.loc[bom_df['foreign_gross'].isna(), 'foreign_gross'] = (
    mean_total - bom_df.loc[bom_df['foreign_gross'].isna(), 'domestic_gross']
)
```

```
In [110]: #Count for the remaining missing values
round((bom_df.isnull().sum()/ len(bom_df))*100,2)
```

```
Out[110]: title          0.0
studio          0.0
domestic_gross  0.0
foreign_gross   0.0
year           0.0
Total_gross     0.0
dtype: float64
```

The remaining data has a few missing values that can be dropped now without having a significant effect on our data

```
In [99]: bom_df= bom_df.dropna()
```

```
In [100]: #Checking the last rows to see if the remaining missing values are all dropped
bom_df.tail()
```

```
Out[100]:
```

| | title | studio | domestic_gross | foreign_gross | year | Total_gross |
|------|---|--------------|----------------|---------------|------|-------------|
| 3275 | I Still See You | LGF | 1400.0 | 1500000.0 | 2018 | 1501400.0 |
| 3286 | The Catcher Was a Spy | IFC | 725000.0 | 229000.0 | 2018 | 954000.0 |
| 3309 | Time Freak | Grindstone | 10000.0 | 256000.0 | 2018 | 266000.0 |
| 3342 | Reign of Judges: Title of Liberty - Concept Short | Darin Southa | 93200.0 | 5200.0 | 2018 | 98400.0 |
| 3353 | Antonio Lopez 1970: Sex Fashion & Disco | FM | 43200.0 | 30000.0 | 2018 | 73200.0 |

The Movies DataBase (TMDB)

Checking any null values in the dataset


```
In [43]: round((tmdb_df.isnull().sum()/ len(tmdb_df))*100,2)
```

```
Out[43]: Unnamed: 0      0.0
genre_ids      0.0
id             0.0
original_language  0.0
original_title  0.0
popularity     0.0
release_date   0.0
title          0.0
vote_average   0.0
vote_count     0.0
dtype: float64
```

There are no null values in the dataset above and everything is as it should be.

The Numbers Dataset

```
In [20]: round((tn_budget.isnull().sum()/ len(tn_budget))*100,2)
```

```
Out[20]: id      0.0
release_date  0.0
movie         0.0
production_budget  0.0
domestic_gross  0.0
worldwide_gross  0.0
dtype: float64
```

This dataset has no missing values but we need to check for info and find out if the datasets are in the correct format

```
In [113]: tn_budget.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 6 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    5782 non-null   int64
 1   release_date          5782 non-null   object
 2   movie                 5782 non-null   object
 3   production_budget     5782 non-null   object
 4   domestic_gross        5782 non-null   object
 5   worldwide_gross       5782 non-null   object
dtypes: int64(1), object(5)
memory usage: 271.2+ KB
```

Some of these columns need to be floats but they are object type.

```
In [21]: #Convert the date column into a pandas-recognizable datetime format

tn_budget['release_date'] = pd.to_datetime(tn_budget['release_date'])
```

```
In [22]: columns = ['domestic_gross', 'production_budget', 'worldwide_gross']

# Remove $ and commas, then convert to float

for col in columns:
    tn_budget[col] = tn_budget[col].replace(r'[\$,]', '', regex=True).astype(float)
```

In [116]: *#Checking if the corrections are in place*
 tn_budget.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 6 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    5782 non-null   int64
 1   release_date          5782 non-null   datetime64[ns]
 2   movie                 5782 non-null   object
 3   production_budget     5782 non-null   float64
 4   domestic_gross        5782 non-null   float64
 5   worldwide_gross       5782 non-null   float64
dtypes: datetime64[ns](1), float64(3), int64(1), object(1)
memory usage: 271.2+ KB
```

From here we then move to data analysis

DATA ANALYSIS

Box Office Mojo

Calculating the trends of both foreign and domestic gross over the years

In [117]: *#Using the Box Office Mojo dataset*
 bom_df.head()

Out[117]:

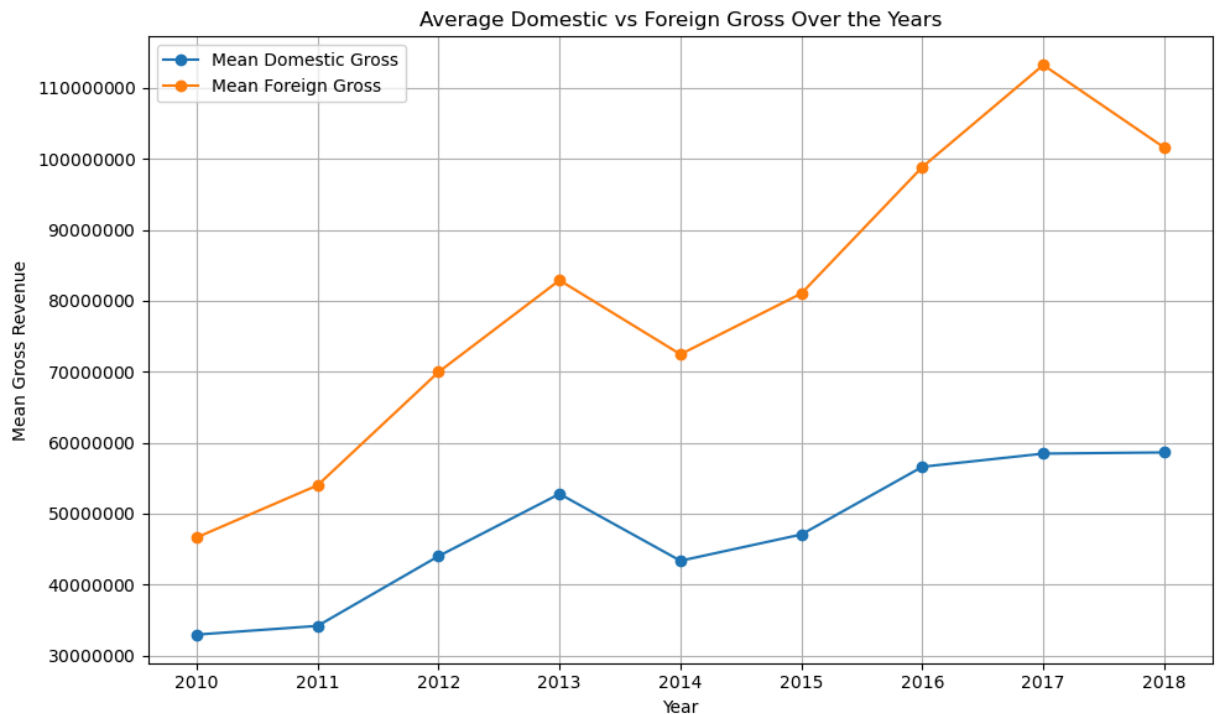
| | title | studio | domestic_gross | foreign_gross | year | Total_gross |
|---|---|--------|----------------|---------------|------|--------------|
| 0 | Toy Story 3 | BV | 415000000.0 | 652000000.0 | 2010 | 1.067000e+09 |
| 1 | Alice in Wonderland (2010) | BV | 334200000.0 | 691300000.0 | 2010 | 1.025500e+09 |
| 2 | Harry Potter and the Deathly Hallows Part 1 | WB | 296000000.0 | 664300000.0 | 2010 | 9.603000e+08 |
| 3 | Inception | WB | 292600000.0 | 535700000.0 | 2010 | 8.283000e+08 |
| 4 | Shrek Forever After | P/DW | 238700000.0 | 513900000.0 | 2010 | 7.526000e+08 |

In [118]: *#Grouping the data*
 gross_by_year = bom_df.groupby('year')[['domestic_gross', 'foreign_gross', 'Total_gross']].mean()

Visualizing comparison between the domestic gross and foreign gross

```
In [119]: #Create plots.
#Use ticker to ensure visualizations are well labelled

import matplotlib.ticker as ticker
plt.figure(figsize=(10, 6))
plt.plot(gross_by_year['year'], gross_by_year['domestic_gross'], label='Mean Domestic Gross', ma
plt.plot(gross_by_year['year'], gross_by_year['foreign_gross'], label='Mean Foreign Gross', mark
plt.gca().yaxis.set_major_formatter(ticker.ScalarFormatter(useMathText=True))
plt.ticklabel_format(style='plain', axis='y') # Ensures full number display
plt.xlabel('Year')
plt.ylabel('Mean Gross Revenue')
plt.title('Average Domestic vs Foreign Gross Over the Years')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



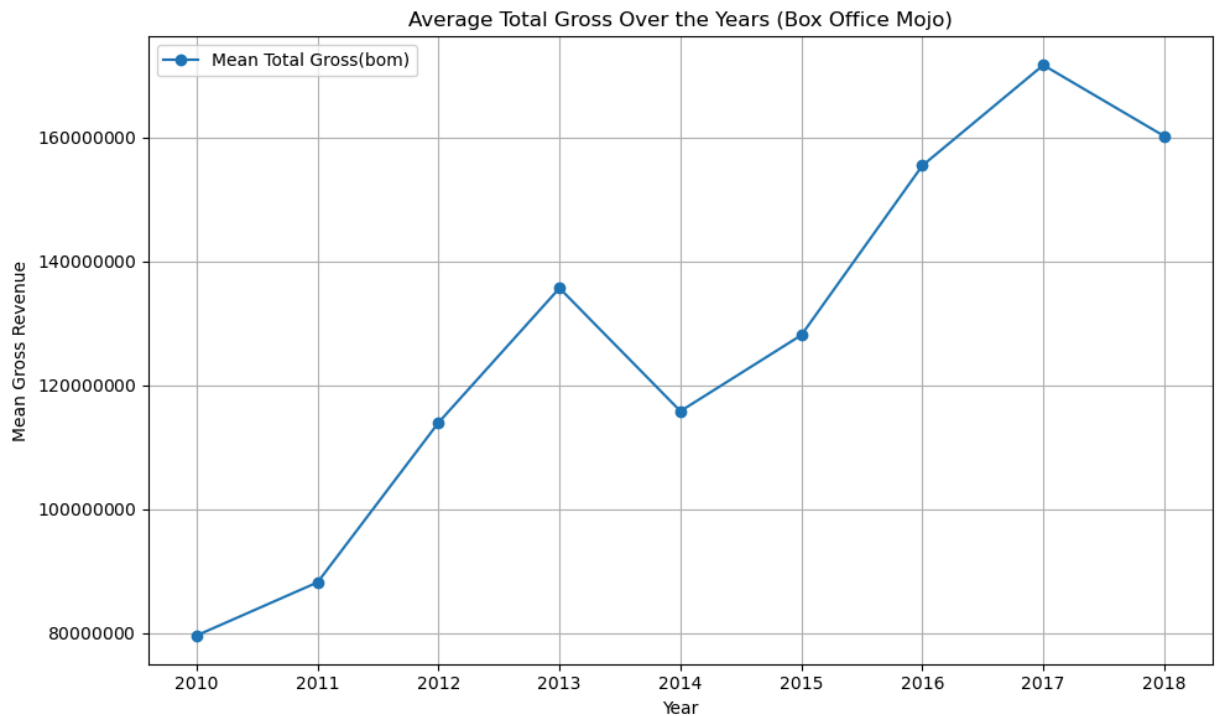
```
In [120]: gross_by_year.value_counts()
```

```
Out[120]: year  domestic_gross  foreign_gross  Total_gross
2010  3.296345e+07  4.665380e+07  7.961725e+07  1
2011  3.419203e+07  5.400989e+07  8.820192e+07  1
2012  4.402413e+07  6.997113e+07  1.139953e+08  1
2013  5.279425e+07  8.290121e+07  1.356955e+08  1
2014  4.336745e+07  7.246439e+07  1.158318e+08  1
2015  4.707417e+07  8.103932e+07  1.281135e+08  1
2016  5.661299e+07  9.886140e+07  1.554744e+08  1
2017  5.847027e+07  1.132326e+08  1.717028e+08  1
2018  5.863697e+07  1.015959e+08  1.602328e+08  1
dtype: int64
```

There is a visible increase in the mean foreign gross from 498,986,300 in 2010 to 1,071,536,000 in 2018

Now visualize The total gross of both domestic and foreign gross

```
In [125]: plt.figure(figsize=(10, 6))
plt.plot(gross_by_year['year'], gross_by_year['Total_gross'], label='Mean Total Gross(bom)', mar
plt.gca().yaxis.set_major_formatter(ticker.ScalarFormatter(useMathText=True))
plt.ticklabel_format(style='plain', axis='y') # Ensures full number display
plt.xlabel('Year')
plt.ylabel('Mean Gross Revenue')
plt.title('Average Total Gross Over the Years (Box Office Mojo)')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



Increase in gross is more visible in the total gross, increasing from 814,415,800 in 2010 to 1,431,640,000 in 2018.

Determining the highest grossing studios

```
In [135]: bom_df['studio'].value_counts()
```

```
Out[135]: Uni.      141
Fox        134
WB         130
Sony       105
BV         102
...
AGF         1
Icar.       1
MPFT        1
KC          1
Darin Southa 1
Name: studio, Length: 172, dtype: int64
```

```
In [137]: #Grouping the data first, by studio and calculate the total gross per category

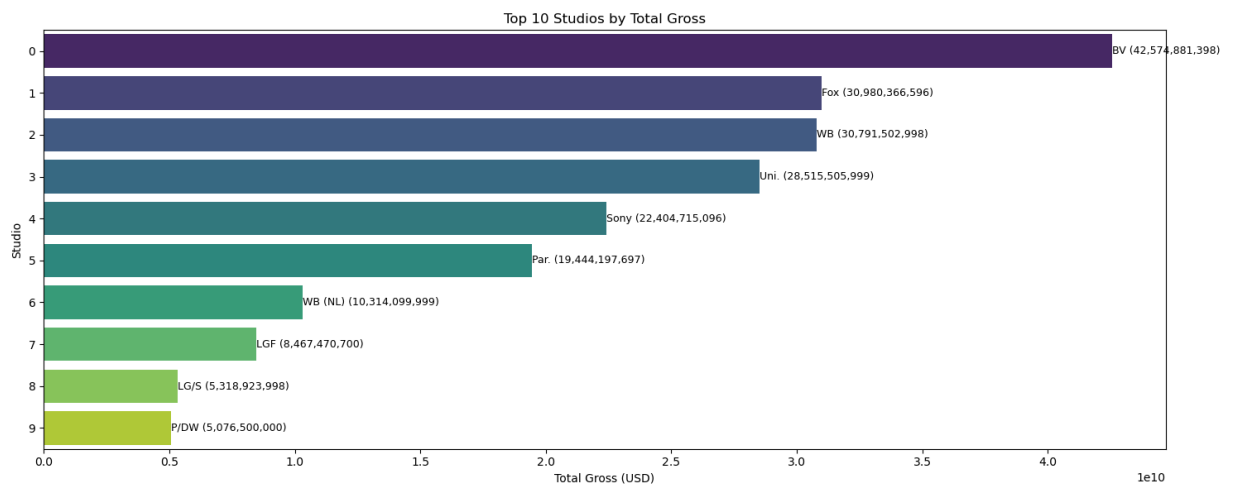
studio_gross = bom_df.groupby('studio')['Total_gross'].sum().sort_values(ascending=False).head(10)

#Plot the visualization
plt.figure(figsize=(15,6))
sns.barplot(x=studio_gross.values, y=studio_gross.index, palette='viridis', )
plt.gca().yaxis.set_major_formatter(ticker.ScalarFormatter(useMathText=True))
plt.ticklabel_format(style='plain', axis='y')

# Ensure that it displays the names

for i, (value, name) in enumerate(zip(studio_gross.values, studio_gross.index)):
    plt.text(value + 1e6, i, f'{name} ({value:,.0f})', va='center', fontsize=9)

plt.title("Top 10 Studios by Total Gross")
plt.xlabel("Total Gross (USD)")
plt.ylabel("Studio")
plt.tight_layout()
plt.show()
```



The top ten most grossing studios are BV down to P/DW with their respective total gross alongside each.

The Numbers Dataset

The numbers production budget over the years

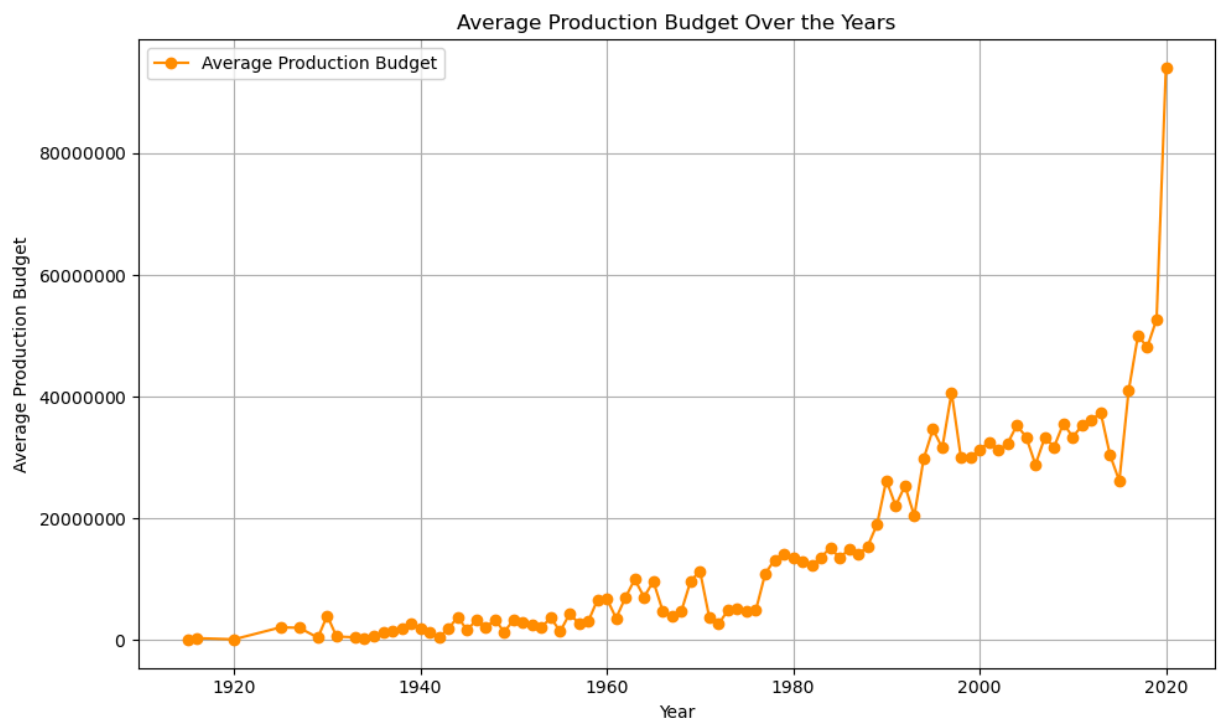
```
In [123]: # Retrieving the year from the date time

tn_budget['year'] = tn_budget['release_date'].dt.year

# Group by year
avg_budget_by_year = tn_budget.groupby('year', as_index=False)['production_budget'].mean()
```

In [126]: # Plotting the production budget

```
plt.figure(figsize=(10, 6))
plt.plot(avg_budget_by_year['year'], avg_budget_by_year['production_budget'], label='Average Pro
plt.gca().yaxis.set_major_formatter(ticker.ScalarFormatter(useMathText=True))
plt.ticklabel_format(style='plain', axis='y')
plt.xlabel('Year')
plt.ylabel('Average Production Budget')
plt.title('Average Production Budget Over the Years')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



We can see an increase from 1,100,000 in 1915 to 940,000,000 in 2020, thus indicating that the budget increased gradually over time.

In [127]: avg_budget_by_year .value_counts()

```
Out[127]: year  production_budget
1915  1.100000e+05      1
1916  2.929535e+05      1
1995  3.463400e+07      1
1994  2.978620e+07      1
1993  2.037973e+07      1
..
1954  3.636667e+06      1
1953  2.128000e+06      1
1952  2.423333e+06      1
1951  2.958333e+06      1
2020  9.400000e+07      1
Length: 96, dtype: int64
```

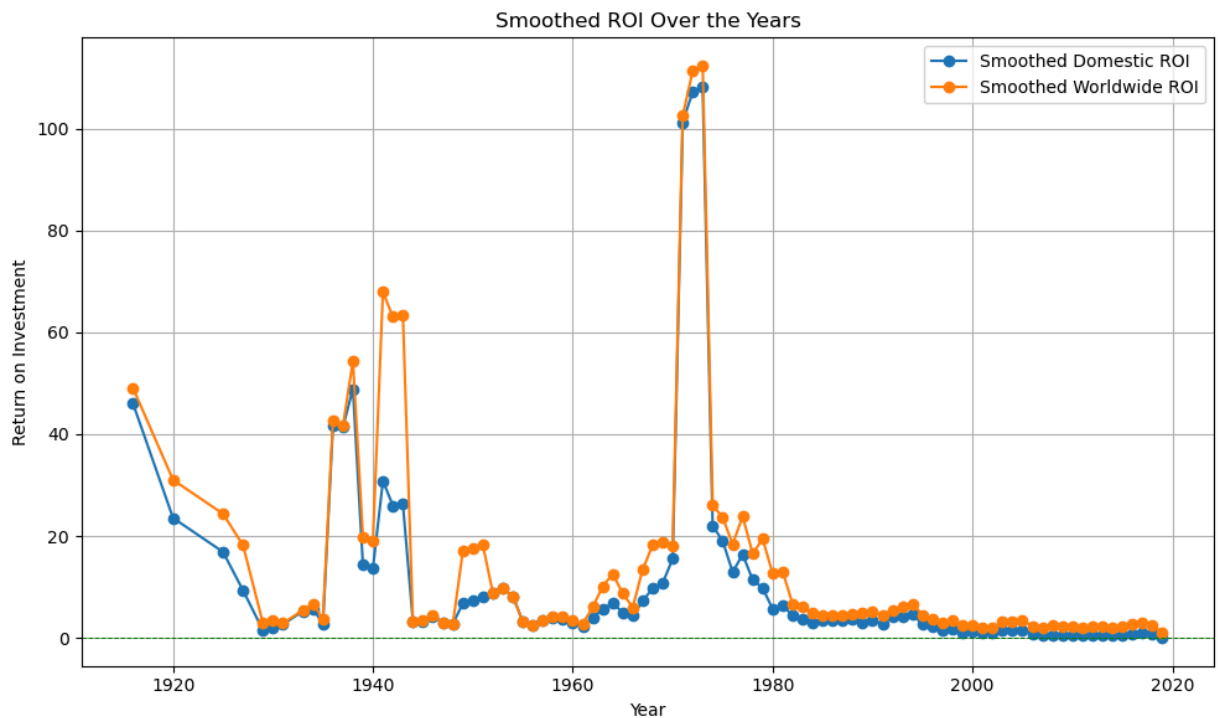
Identifying trends of ROI over the years

```
In [128]: tn_budget['domestic_roi'] = (tn_budget['domestic_gross'] - tn_budget['production_budget']) / tn_
tn_budget['worldwide_roi'] = (tn_budget['worldwide_gross'] - tn_budget['production_budget']) / t
```

```
In [129]: tn_budget['year'] = tn_budget['release_date'].dt.year
roi_by_year = tn_budget.groupby('year', as_index=False)[['domestic_roi', 'worldwide_roi']].mean()
```

```
In [134]: roi_by_year['domestic_roi_smooth'] = roi_by_year['domestic_roi'].rolling(window=3, center=True).
roi_by_year['worldwide_roi_smooth'] = roi_by_year['worldwide_roi'].rolling(window=3, center=True)

# Plotting smoothed curves
plt.figure(figsize=(10, 6))
plt.plot(roi_by_year['year'], roi_by_year['domestic_roi_smooth'], label='Smoothed Domestic ROI',
plt.plot(roi_by_year['year'], roi_by_year['worldwide_roi_smooth'], label='Smoothed Worldwide ROI')
plt.axhline(0, color='green', linestyle='--', linewidth=0.7)
plt.xlabel('Year')
plt.ylabel('Return on Investment')
plt.title('Smoothed ROI Over the Years')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



The curve is very unpredictable in both the domestic and worldwide ROI. In the 15s to 70s (years), the ROI was high reaching over 50%, however in the recent years this has gone down to less than 10% annually.

The Movies Database (TMDB)

Movie language and popularity

```
In [44]: tmdb_df['original_language'].value_counts()
```

```
Out[44]: en    23291
fr      507
es      455
ru      298
ja      265
...
bo         1
si         1
sl         1
hz         1
dz         1
Name: original_language, Length: 76, dtype: int64
```

At the top is the English movie with 23,291.

RECOMMENDATIONS

Partnering with High-Grossing Studios such as Sony Pictures, Warner Bros, and 20th Century Fox have shown immense-grossing films and Collaborating or buying their ideas of production and distribution strategies could enhance commercial outcomes.

Targeting International Markets; International releases show greater predictability and performance thus expanding production efforts to cater to global audiences can enhance market penetration and revenue stability.

Maximizing Production Budgets Strategically; Investing in quality production, casting, and marketing is essential to securing long-term financial returns.

Investing in High-Demand Language Markets; Strategic investment in languages such as English, French, Spanish, and Russian that consistently demonstrate higher popularity and commercial success is advised to capture broader audience bases.

In []: